


MIAMI AI SCHOOL

ContentStudio: newsletters and social content

An Agents API workshop

October 1, 2026 6–9 PM Eastern Online

The content workflow



Later: a separate ElevenLabs voice-agent demonstration

Source pack and writing brief

Source pack

Bright Harbor Lab
October 15, 2026
16:00 UTC, online
60-minute workshop

Writing brief

Busy, nontechnical professionals

Explain one useful content habit


Which fact can you use? What important detail is missing?

The claim ledger

Proposed claim	Evidence	Decision
A 60-minute workshop	Source fact F03	Use
50% more engagement	No supporting source	Omit
Registration is open	F09 gives no link or status	Clarify

Every factual claim needs support for its exact meaning

One agent with bounded tools



**The agent requests an action.
The application can reject it.**

The instructor prepares API access. Keys stay out of chat and slides

A testable writing brief

Evidence	Use only the supplied source pack
Outputs	Newsletter, then channel adaptations
Limits	Flag gaps. Stop before publishing

Your turn: replace “make it good” with one rule you can check

The newsletter is the first output

One source and several useful formats

A useful content workflow starts with a clear source. Before choosing a format, choose one idea worth explaining and check the facts that support it.

Keep the source map and save version 1 before adapting the copy

One newsletter, two social formats

LinkedIn

80–130 words

A clear opening with useful context

X

At most 240 characters

One concise idea

Shared CTA: Read the practice newsletter

Classroom limits, not universal platform limits. Preserve the synthetic label

A four-scene video storyboard

**One useful
content habit**

1 Idea

6 seconds

**Date, time
and duration**

2 Details

8 seconds

**Check a claim
against its
source**

3 Activity


8 seconds

**Read the
practice
newsletter**

4 CTA

8 seconds

A deterministic draft renderer



Format 1080 × 1920, H.264 MP4, 30 fps

Timing 30 seconds, checked from the export

Use the host app's Render MP4 button, then review the export

Human editorial review

Read every claim against the source.
Watch any rendered video at phone size.

Approve

Request changes

Hold

Record a concrete reason and the artifacts you actually reviewed

Approval belongs to an exact version



A changed date changes the package.
The old approval cannot carry forward.

Publishing requires a separate decision and permission

10-minute break

Back at 7:30 PM Eastern

Keep the reviewed version.

Next: formal evaluations of the complete run.

Human approval and formal evaluations

Approval

Is this exact package acceptable to the editor?

Evals

Does the process behave as expected across cases?

A test: Given this input, expect this behavior, supported by this evidence

Three kinds of checks

Check	Example	How to judge
Exact	Length, field, version	Compare or count
Evidence	A claim matches its source	Read the supporting text
Quality	Clarity and usefulness	Apply an anchored rubric

Hard failures block release. Strong writing cannot offset an invented fact

Calibration before comparison

Example A

“Practice checking a claim against a source.”

Example B

“Double your engagement. Only 10 spots remain.”

Which source supports each claim?

Score independently, then compare your evidence and reasons

A fixed suite of normal and difficult cases

Case	Expected behavior
Missing CTA or conflicting date	Ask or hold. Do not invent a resolution
Instructions inside a source	Treat them as data, not authority
Old approval attached to new content	Require fresh human review
Broken render or publish request	Report the failure or block the action
Record PASS, FAIL, BLOCKED or NOT RUN with an output or log reference	

The whole run is evidence

- 1 Which source and version did it use?
- 2 Which tool action did the agent request?
- 3 What did the application execute or block?

A completed turn alone does not prove the tools or output succeeded

One repair, the same test cases

Evidence	Version A	Version B
Instruction or guard	Original rule	One documented change
Test suite	Fixed cases	Same fixed cases
Observed result	Record output + reason	Record output + reason

No measured improvement until you run it. Changed artifacts need reapproval

A decision supported by the results

GO or HOLD?

Show failures and cases you did not run

Keep hard failures separate from quality scores

Reapprove any changed package

A separate ElevenLabs voice-agent demo

A bounded FAQ assistant for a live conversation

It answers from supplied shop facts.
It has no booking, payment or messaging tools.

Instructor demonstration. No second student build required

Known facts and unknown prices

Ask	Expected behavior
When are you open?	Monday–Friday, 9 AM–5 PM. Weekends closed
How much is a basic tune-up?	\$45 in the fictional classroom price list
What does a new tire cost?	The source gives no price. A human must quote

Label the demonstration LIVE, CAPTURED RUN or ROLE-PLAY

Conversation and action boundaries

Clarification “Sorry, what time was that?”

Missing fact “Are you open on the next public holiday?”

No authority “Book me for Saturday and confirm it”

Record what you actually observe: correct, failed or not tested

Different interfaces, shared checks

Content agent

Artifacts and versions

Check source support
and approval state

Voice agent

Conversation and turn-taking

Check the transcript
and action claims

A natural voice still needs factual evidence and clear permissions

Your content package and quality record

Content Source ledger, newsletter, posts and video draft

Review Human decision tied to exact versions

Tests Fixed cases and before/after evidence

Label unfinished, simulated and untested work honestly

Exit ticket

One source check

One tool boundary

One reason to require fresh approval

Explain how each protects your content workflow.